

Dr. Babasaheb Ambedkar Technological University, Lonere, Rigad

A PROJECT REPORT ON

“AI for Legal Case Outcome Prediction”

Project ID: (CSE-AI-06/2025-26)

For the Partial Fulfillment

**Bachelor of Technology in Department of Computer Science Engineering
(Artificial Intelligence)**

SUBMITTED BY:

Amay Lohar	2264191913014
Nikhil Raghuwanshi	2264191913036
Shital Kalekar	2264191913063
Shraddha Dhanwate	23064191913509

UNDER THE GUIDANCE OF

Mrs. Priyanka Patil



Department Of Computer Science Engineering (Artificial Intelligence)

PCET- NMVPM's,

Nutan College of Engineering and Research, Talegaon, Pune

“Vision of CSE – Artificial Intelligence”

“Excellence in the domain of Artificial Intelligence towards serving the greater cause of society and globally recognized for contributing professional engineers with a commitment to readiness of industry-oriented skill through potential research.”

2025-2026



Nutan Maharashtra Vidya Prasarak Mandal's
NUTAN COLLEGE OF ENGINEERING AND RESEARCH
Under Administrative Support of Pimpri Chinchwad
Education Trust



CERTIFICATE

This is to certify that the project report entitled “**AI for Legal Case Outcome Prediction**”, which is being submitted by, **Amay Lohar, Nikhil Raghuwanshi, Shital Kalekar and Shraddha Dhanwate** as partial fulfillment for the Degree of Bachelor of Technology (Computer Science Engineering –Artificial Intelligence) of **DBATU, Lonere**

This is bonafide work carried out under my supervision and guidance.

Place: Pune

Date:

Mrs. Priyanka Patil
Project Guide

Mrs. Priyanka Vyas
Project Coordinator

Dr. Prashant Ahire
Head of Department

Dr. Pramod Patil
Director

External Examiner

SEAL

ACKNOWLEDGEMENT

The satisfaction & euphoria that accompany the successful completion of any task would be incomplete without the mention of people who made it possible. So we acknowledge all those whose guidance and encouragement served as a beacon light & crowned my efforts with success.

We have immense pleasure in expressing thanks to the director Dr. Pramod Patil for providing all the facilities for the successful completion of the project. With due respect, we thank my Head of Department Dr. Prashant Ahire Department of Computer Science Engineering (Artificial Intelligence), for his motivating support, keen interest which kept spirits alive all through.

We would like to express thanks to my guide **Mrs. Priyanka Patil**, Department of Computer Science Engineering (Artificial Intelligence) who has guided us throughout the completion of this project.

Finally we would like to thank all the teaching and non-teaching staff and all our friends who have rendered their support in the completion of this report.

Name and Signature of student:

Sr.no	Name	Sign
1	Amay Lohar	
2	Nikhil Raghuwanshi	
3	Shital Kalekar	
4	Shraddha Dhanwate	

ABSTRACT

The Indian legal system handles an large volume of cases annually, with judgments often covering hundreds of pages. Consequently, physically analysing precedents and forecasting outcomes is highly complex and time-consuming. This paper presents a production-ready Artificial Intelligence (AI) framework designed to predict Indian Supreme Court case outcomes however addressing modern legal transitions. Our novel architecture utilizes InLegalBERT for deep appropriate surrounding, augmented by a dynamic translation layer that maps legacy Indian Penal Code (IPC) references to the new Bharatiya Nyaya Sanhita (BNS) using structured JSON dictionaries. To overcome the standard 512-token limit of transformers, we fixed a structure-aware chunking pipeline that isolates 'Facts' and 'Arguments' using targeted Regular Expressions. Additionally, to satisfy the essential need for judicial transparency, we integrated SHAP(SHapley Additive exPlanations) with spaCy Named Entity Recognition (NER) to sort out grammatical noise, providing specialists with highly focused, interpretable decision drivers. The entire predictive engine is deployed as a secure SaaS platform utilizing a Streamlit frontend and a Supabase (PostgreSQL) backend for encoded user authentication and case history management. Our results show that this end-to-end system accurately predicts case outcomes and offers brief, explainable justifications, making it highly feasible for real-world application in the Indian legal ecosystem.

Index Terms—Predicting the outcome of legal cases, legal reasoning, the Indian legal system, text classification, explainable artificial intelligence, legal document analysis, court judgment analysis, criminal-judicial data analysis, and legal decision support systems

Keywords: *Legal Case Outcome Prediction, Indian Legal System, Artificial Intelligence (AI), Text Classification, Explainable AI (XAI), Legal Document Analysis, Court Judgment Analysis, Criminal-Judicial Data Analysis, Decision Support Systems, InLegalBERT, SHAP Explainability, Named Entity Recognition (NER), IPC to BNS Mapping, SaaS Legal Platform.*

INDEX

Sl. No.	Content	Page No.
	Certificate	2
	Acknowledgment	3
	Abstract	4
1.	Introduction	8
	1.1 History	9
	1.2 Literature Review	10
	1.3 Gap Findings	11
	1.4 Problem Formulation	12
	1.5 Problem Solution	13
	1.6 Objectives	13
2.	System Architecture and Methodology	14
	2.1 Block Diagram of Proposed Work	16
	2.2 Methodology	17
	2.3 Algorithm	18
	2.4 Flowchart	20
3	Experimental Results	21
4	Conclusions	23
	Research Papers Published/Certificates	24
	References	25

LIST OF FIGURES

Figure No.	Caption	Page No.
2	Architecture of the Legal Justice Prediction System	15
2.1	The Input Interface for the Legal AI System	17
2.4	Detailed Case Analysis with Predictive Explanations	20
3	Model accuracy comparison	21
3.2	Confusion matrix	22
3.3	Feature importance graph based on SHAP values	23
3.4	ROC curve of the InLegalBERT classifier	23
3.5	Feature correlation heatmap	24
3.6	Algorithm runtime comparison	24
3.7	Prediction Confidence Distribution (PCD)	25
3.8	The platform login interface for secure access	25
3.9	The Secure Case History dashboard	26
3.10	The detailed case analysis dashboard	26

LIST OF TABLES

Table No.	Caption	Page No.
1.2	Literature Survey	10
3	Metrics Achieved on Unseen Data	21

1. Introduction

The Indian legal system oversees a vast volume of cases annually, with judicial verdicts frequently extending across hundreds of pages of dense documentation. This sheer volume makes the manual analysis of historical precedents and the forecasting of future case outcomes an exceptionally complex, labor-intensive, and time-consuming process for legal practitioners. Consequently, there is an urgent industry requirement for advanced computational tools to assist in navigating this complexity.

This project presents a production-ready Artificial Intelligence (AI) framework specifically engineered to predict the outcomes of Indian Supreme Court cases while proactively addressing modern transitions in the legal landscape. Predicting the resolution of a court case is a critical task within legal analytics, serving to foresee judicial decisions by synthesizing written laws and historical precedents. Predictive methodologies in this field are generally classified into three distinct types: the **socio-political approach**, which examines extrajudicial factors like a judge's background or ideological inclinations; the **linguistic approach**, which identifies patterns in document structure, sentiment, and vocabulary; and the **legal-centric approach**, which focuses on the specific facts of each case alongside relevant statutory laws.

A significant computational hurdle identified in current Indian legal tech is the "temporal shift"—the evolution of legal interpretations and societal norms over time, which often renders older models obsolete. Furthermore, the recent legislative transition from the Indian Penal Code (IPC) to the Bharatiya Nyaya Sanhita (BNS) has created a unique computational challenge, as legacy AI models were trained exclusively on outdated legal terminology. Our system addresses this by introducing a dynamic middleware layer that automatically translates historical laws into updated forms before inference, ensuring that the model maintains high relevance to current legal standards.

Finally, to bridge the divide between academic research and commercial viability, the entire workflow is integrated into a secure, full-stack architecture. This framework utilizes Supabase for robust lawyer authentication and cloud-based case storage, demonstrating the model's readiness for practical, real-world deployment in the Indian judicial ecosystem.

1.1 History

The field of legal predictive analytics has evolved significantly, transitioning from primitive manual assessment methods to advanced computational frameworks. Historically, predictive modeling in the legal domain has been categorized into three distinct methodologies: the socio-political approach, which emphasizes extrajudicial factors such as a judge's personal background and

ideological inclinations ; the linguistic approach, which identifies trends in document structure, sentiment, and word frequency within legal texts ; and the legal-centric approach, which concentrates specifically on case facts and the application of related statutory laws.

Over the last decade, computational legal technology has advanced from traditional classification algorithms—such as Support Vector Machines (SVM) and Random Forest models—toward sophisticated, domain-specific Natural Language Processing (NLP) architectures. Early benchmarks, such as those established by Katz et al. (2017), focused on modeling judicial behavior over long horizons, while more recent advancements have leveraged Transformer-based models like BERT to handle the semantic density of legal language. The field has progressively shifted toward addressing the complexities of "temporal shifts," where legal interpretations evolve alongside societal norms, necessitating models that can adapt to changing statutory structures. The current trajectory of legal history in technology is defined by the integration of Explainable Artificial Intelligence (XAI), moving away from "black-box" systems toward transparent frameworks that provide justifiable reasoning pathways, ensuring that automated predictions meet the rigorous accountability standards of the Indian judicial ecosystem.

1.2 Literature Review

The literature review on Indian legal technology highlights an increasing requirement for multilingual understanding and interpretable reasoning to bridge the gap between complex statutory frameworks and practical law enforcement systems. Recent improvements have focused on domain-specific benchmarks, such as IL-TUR, which have expanded research across many Indian languages and allowed for task-specific evaluation of legal NLP systems. However, significant challenges persist in cross-lingual reasoning and in optimizing the baseline performance of Large Language Models (LLMs) when utilized for Indian legal texts, which are highly complex in both structure and meaning.

Existing literature notes a shortage of large-scale annotated corpora and a lack of representation for regional languages. A primary technical hurdle is that Indian legal documents frequently exceed 4,000 words, far surpassing the 512-token limitation of traditional transformer architectures such as BERT. This limitation necessitates the development of advanced chunking and large-document modeling strategies. Furthermore, the high semantic density of statutory interpretation makes it difficult for black-box neural models to create transparent outputs that satisfy judicial accountability standards.

To address these boundaries, current research has pivoted toward Explainable Artificial Intelligence (XAI) frameworks. While XAI-based models may display slightly condensed predictive precision compared to fully optimized black-box systems, they significantly enhance user trust by providing

justifiable reasoning pathways. The urgency for such systems has intensified with India's legislative change from the Indian Penal Code (IPC) to the Bharatiya Nyaya Sanhita (BNS), creating a critical need for adaptable computational frameworks capable of handling growing statutory structures. Current benchmarks in the field are illustrated in the comparative analysis below:

Sr. No.	Author	Dataset	Method	Contribution
1	Katz et al., 2017 [21]	US Supreme Court	Random Forest	Long-horizon behaviour modelling
2	Malik et al., 2021 [23]	ILDC Corpus	Hierarchical model	Indian legal prediction benchmark
3	Liu et al., 2024 [24]	CAIL2018	LegalDuet	Dual-view reasoning
4	Nigam et al., 2024 [27]	NyayaAnumana	Legal LLM	Large-scale Indian legal baseline
5	Garzo et al., 2025 [8]	Criminal data	Transformer	Interpretable judicial prediction
6	LIPE Model, 2025 [18]	US Supreme Court	Ensemble + NLP	Hybrid feature engineering

1.3 Gap Findings

A. Architectural & Computational Gaps

- **Limited Context Understanding:** Traditional machine learning models (such as SVM or Naive Bayes) fail to capture the nuanced legal reasoning and context required to interpret complex judgments.
- **Handling Long Documents:** Most standard Transformer-based models operate under a 512-token limit, leading to the loss of critical information when processing lengthy legal case documents.
- **Poor Handling of Legal Structure:** Existing systems often fail to properly separate specific

document sections like "Facts," "Arguments," and "Judgment," which results in inaccurate predictive modeling.

- **Data Limitations:** Research is frequently constrained by a shortage of large, well-annotated datasets, further exacerbated by the presence of messy, unstructured legal text.

B. Legal Transparency & Adaptability Gaps

- **Lack of Explainability (Black Box Problem):** Many high-accuracy AI models function as "black boxes," providing predictions without transparent justifications, which reduces trust among legal practitioners and fails to meet judicial accountability standards.
- **Static Models:** Existing systems lack adaptability to evolving laws, such as the transition from the Indian Penal Code (IPC) to the Bharatiya Nyaya Sanhita (BNS), leaving them unable to provide relevant or current guidance.

C. Practical Deployment & Linguistic Gaps

- **Multilingual Challenges:** While Indian legal data is inherently multilingual, the majority of existing models are trained primarily on English, creating a barrier to broader applicability.
- **No Real-World Deployment Focus:** A significant portion of current research remains purely theoretical, failing to provide the secure, user-friendly SaaS applications required for day-to-day legal workflows.

These findings confirm that while progress has been made in predictive modeling, there is a lack of end-to-end, explainable, and adaptable systems. The **Legal AI** bridges these gaps by implementing structure-aware chunking, dynamic IPC-to-BNS translation, and SHAP-based explainability within a secure, SaaS-enabled architecture.

1.4 Problem Formulation

The Legal Case Outcome Prediction project aims to develop a robust machine learning system capable of predicting the outcomes of court cases by analyzing historical legal verdicts. In common law systems like India, Supreme Court and High Court judgments serve as the backbone of legal understanding. However, the current manual process of reviewing large volumes of unstructured legal documents to identify relevant precedents and reasoning patterns is inherently time-consuming, subjective, and complex.

The fundamental problem lies in the disconnect between the massive influx of legal data and the lack of automated, explainable tools to process it. Specifically, our formulation focuses on three key technical and structural challenges:

- **Extraction of Structured Evidence:** Converting unstructured, verbose legal texts into actionable data points while maintaining the integrity of the original arguments.
- **Adaptability to Legislative Evolution:** Ensuring that the model remains functional amidst

significant statutory transitions—specifically, the shift from the Indian Penal Code (IPC) to the Bharatiya Nyaya Sanhita (BNS)—which renders legacy models trained on outdated terminology ineffective.

- **Balancing Accuracy and Transparency:** Addressing the requirement for models to provide not just a binary outcome (e.g., "Dismissed" or "Allowed") but also interpretable justifications, as judicial accountability requires transparency in AI-assisted decision-making.

By training machine learning models on high-quality, historical case data, the system aims to automate the prediction of outcomes, such as whether an appeal will be allowed or dismissed. The core objective is to reduce the burden of manual legal research, enhance consistency in judicial analysis, and provide data-driven insights through a secure, scalable SaaS platform. This formulation transforms the manual task of judgment review into a streamlined, automated pipeline, assisting legal professionals and researchers in navigating Indian jurisprudence more capably.

1.5 Problem Solution

To address the inherent complexities of the Indian judicial system and the limitations of existing predictive models, we have developed an end-to-end, explainable AI framework designed for real-world legal decision support. The solution is formulated around a modular, full-stack architecture that transforms unstructured legal texts into structured, actionable intelligence.

Our solution is built upon four foundational pillars:

- **Advanced NLP & Structural Chunking:** To overcome the 512-token limit of standard Transformer architectures, we implemented a structure-aware chunking pipeline. This logic utilizes targeted Regular Expressions to strictly isolate the "Facts" and "Arguments" sections of legal judgments, ensuring that the inference engine processes the most legally dense and unbiased information.
- **Dynamic Statutory Translation (IPC-to-BNS Middleware):** To ensure the model remains relevant amidst legislative transitions, we integrated a vibrant JSON-based transformation layer. This middleware dynamically maps legacy Indian Penal Code (IPC) references to updated Bharatiya Nyaya Sanhita (BNS) provisions, allowing the model to evaluate historical cases using contemporary legal standards.
- **Explainable AI (XAI) Integration:** To solve the "black-box" problem, we integrated SHAP (SHapley Additive exPlanations) coupled with a spaCy Named Entity Recognition (NER) filter. This ensures that the system does not just provide a binary verdict, but highlights specific legal entities, laws, and case-specific decision drivers, providing practitioners with a transparent and verifiable decision path.

- **Secure SaaS Implementation:** The entire engine is wrapped in a production-ready, secure SaaS platform. By utilizing a Streamlit frontend and a Supabase (PostgreSQL) backend, we provide lawyers with robust, encrypted user authentication, a persistent cloud vault for case history, and an intuitive dashboard that enhances judicial consistency and saves research time.

Through synergising domain-specific transformer models—specifically fine-tuned InLegalBERT—with robust backend infrastructure, this solution effectively mitigates the risks of manual research and data loss. The platform provides data-driven insights that are grounded in actual case facts, empowering legal professionals and researchers to navigate Indian jurisprudence with greater speed, accuracy, and interpretability. This integrated approach ensures that the Legal AI Hub is not merely a research prototype, but a scalable, commercial-grade tool ready for integration into the Indian legal ecosystem.

1.6 Objectives

- **Data Acquisition and Management:** To systematically collect historical case judgments from the Supreme Court of India and interconnected legal repositories to create a reliable training corpus.
- **Intelligent Data Preprocessing:** To clean and structure unstructured legal texts by dynamically extracting essential components such as "Facts" and "Arguments," while implementing rigorous filtering to prevent data leakage from final judgment sections.
- **Statutory Evolution Mapping:** To implement a dynamic JSON-based transformation middleware capable of mapping legacy Indian Penal Code (IPC) references to the contemporary Bharatiya Nyaya Sanhita (BNS) framework, ensuring the model remains legally relevant.
- **Semantic Entity Recognition:** To apply advanced Natural Language Processing (NLP) techniques, augmented by spaCy Named Entity Recognition (NER), to isolate critical legal keywords and eliminate grammatical noise.
- **Predictive Model Development:** To fine-tune and deploy transformer-based architecture (InLegalBERT) to perform high-precision classification of case outcomes.
- **Secure Cloud Infrastructure:** To implement a robust, scalable cloud backend utilizing Supabase for encrypted user authentication, secure role-based access, and persistent historical case tracking.
- **Explainable User Interface:** To develop a professional-grade, interactive web interface using Streamlit that empowers legal professionals to analyze cases rapidly, supplemented by explainable visual decision drivers generated via SHAP (SHapley Additive exPlanations).

2. System Architecture and Methodology

The proposed system utilizes a modular, full-stack architecture to ensure scalability, security, and high-performance predictive analysis. As illustrated in the architecture diagram, the workflow is designed as a continuous pipeline that bridges the gap between raw legal documents and interpretable judicial forecasts.

Key Components of the Architecture:

- **Legal Professional (User):** The primary interface actor who interacts with the system through a secure, state-based web portal.
- **Streamlit Frontend (State-Based UI):** Acts as the interactive entry point, managing user sessions and routing requests across various modules (Prediction, Vault, History) while ensuring a responsive user experience.
- **Supabase Backend (Auth & PostgreSQL):** Serves as the robust cloud foundation, providing encrypted user authentication, row-level security, and a persistent vault for historical case data and user logs.
- **Data Preprocessing Engine:** A dedicated module comprising three critical sub-stages:
 - **PyMuPDF (Text Extraction):** Extracts structured text from unstructured PDF legal documents.
 - **JSON Middleware (IPC to BNS Translation):** A novel novelty layer that dynamically maps legacy Indian Penal Code (IPC) references to the contemporary Bharatiya Nyaya Sanhita (BNS) framework.
 - **Structure-Aware Chunking (Facts & Arguments):** Isolates the most legally dense sections of the document to optimize inference and prevent information loss due to standard token limitations.
- **InLegalBERT (PyTorch Inference):** The core transformer-based engine that processes pre-processed text to generate probability distributions for case outcomes.
- **XAI Module (SHAP + spaCy NER):** The final stage which provides transparency. It filters grammatical noise via spaCy Named Entity Recognition and uses SHAP to highlight critical legal tokens that drive the model's predictive decision, providing practitioners with verifiable evidence.




Fig 1. Architecture of the Legal Justice Prediction System

2.1 Block Diagram of Proposed Work

The proposed workflow, visualized in the input interface diagram (Fig 4), delineates the modular navigation and operational flow of the Legal AI system. This flow ensures a seamless experience for legal practitioners, maintaining data integrity from the initial login to the archival of predictive insights.

Workflow Stages:

- **Secure Authentication (Login/Sign-up Portal):** The entry point of the system, providing encrypted access to the platform.
- **Main Dashboard Navbar (State-Based Routing):** A central navigation hub that utilizes state-based routing to provide a Single Page Application (SPA) experience, allowing users to toggle between different analytical modules without page reloads.
- **Document Upload Module (PDF/TXT Dropzone):** A secure ingestion point that enables users to upload legal documents for analysis.
- **Prediction Metric Cards (Jurisdiction & Outcome):** The primary analytical output where the InLegalBERT model displays real-time jurisdiction classification and outcome probability forecasts.
- **XAI & Structure Tabs (SHAP Visuals & Extracts):** An interactive section that visualizes explainable AI (XAI) insights and dynamically extracts case-specific facts and arguments for auditability.
- **Secure Case History (Cloud Vault Database):** The final stage of the workflow, where predicted case data is pushed to a persistent Supabase PostgreSQL cloud vault for encrypted storage and future audit tracking.




Fig 2. The Input Interface for the Legal AI System

2.2 Methodology

The methodology of the **Legal AI** is designed as a modular, full-stack pipeline that facilitates the end-to-end processing of legal judgments. The workflow begins with the ingestion of raw documents and concludes with the generation of explainable, data-driven insights.

1. Data Ingestion and Text Extraction

The initial step involves obtaining legal judgments, which are primarily sourced in PDF format from public repositories such as the Supreme Court of India website. We utilize **PyMuPDF (fitz)** to perform reliable, well-organized text extraction, ensuring that the raw content is ready for secondary analysis.

2. Advanced Preprocessing

The extracted raw text undergoes rigorous cleaning to ensure high-quality model inputs:

- **Text Cleaning:** Regular Expressions (Regex) are employed to remove non-essential elements such as boilerplate text, headers, footers, and page numbers.
- **Dynamic Statutory Translation:** A core novelty of our system is the automated mapping of legacy IPC sections to modern BNS sections using an external JSON dictionary and Regex word boundaries. This ensures the AI evaluates historical cases according to current legal standards.
- **Structure-Aware Document Chunking:** To circumvent the 512-token limit of Transformer models, we employ targeted Regular Expressions to isolate the "FACTS" and "ARGUMENTS" sections. This ranking process ensures that the inference engine prioritizes the most legally dense and unbiased information.

3. Model Inference (InLegalBERT)

The processed text chunks are tokenized and passed through the fine-tuned **InLegalBERT** model. The model utilizes an encoder to generate contextual embeddings, which are then processed by a classification head to generate logits. These logits are converted into final outcome probabilities using the Softmax function:

$$P(y = i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

4. Explainable AI (XAI) Pipeline

To ensure judicial transparency, we integrate **SHAP** (SHapley Additive exPlanations) to assign importance values to each token. These values are piped through a **spaCy Named Entity Recognition (NER)** filter, which assigns a zero-weight to grammatical noise and structural words. This allows the dashboard to exclusively highlight high-effect legal entities and provisions (e.g., "bail," "dismissed," "Section 103 BNS") in red and blue, providing a verifiable decision path.

2.3 Algorithm

The classification involves Linear Model Formulation ($z = w^T x + b$) and a Softmax

adaptation ($P(y = i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$) applied to the embeddings generated by the InLegalBERT encoder.

2.4 Flowchart

1. The explainability pipeline represents a multi-stage workflow designed to convert model logits into human-interpretable decision drivers. This process ensures that judicial forecasting is not only accurate but also transparent and verifiable.
2. Stages of the Pipeline:
3. Cleaned Feature Vectors (Facts & Arguments): The workflow begins with structured input data, where legal documents have been chunked to isolate the most relevant argumentative components.
4. InLegalBERT Prediction Model: The fine-tuned transformer model processes these feature vectors to compute raw logits and outcome probabilities.
5. Explainable AI (XAI) Pipeline: This central module performs two critical functions:
 - a. SHAP Explainer: Utilizes game theory to assign an importance value (Shapley value) to every token, quantifying its contribution to the final prediction.
 - b. spaCy NER Filter: Acts as a specialized legal filter that isolates important legal entities (e.g., specific laws or consequences) while assigning zero-weight to grammatical noise.
6. Interactive UI Visualization: The final output is rendered in the Streamlit dashboard, highlighting critical legal words with red and blue markers. This visual feedback allows legal practitioners to immediately identify the core drivers behind the AI's forecast.




Fig 3. Detailed Case Analysis with Predictive Explanations

3. Experimental Results

The accuracy comparison chart illustrates that our optimized **InLegalBERT** architecture significantly outperforms established models like SVM, Random Forest, and Bi-LSTM. This performance verification confirms that domain-specific transformer models are better equipped to handle the intricacies of Indian court judgments compared to traditional classification algorithms.

The quantitative metrics achieved on the unseen test set are summarized below:

Sr. No.	Metric	Value
1	Accuracy	0.852
2	Precision	0.841
3	Recall	0.835
4	F1-Score	0.837



Fig 4. Model accuracy comparison

These metrics validate high analytical trustworthiness across all primary classification categories. With an accuracy of 85.2% and a balanced F1-score of 0.837, the model demonstrates its efficiency in managing the linguistic complexities of judicial documents.

3.2 Confusion Matrix Analysis

The confusion matrix (Fig 7) visualizes the distribution of actual versus predicted outcomes. The concentration of values along the primary diagonal (428 True Positives and 419 True Negatives) underscores the model's high precision in identifying both "Allowed" and "Dismissed" decisions. By maintaining a low error rate, the system proves its robustness for real-world legal forecasting, ensuring practitioners can rely on the AI's predictive consistency.

	Pred Pos	Pred Neg
Act Pos	428 True Positive	72 False Negative
Act Neg	81 False Positive	419 True Negative

Fig 5. Confusion matrix

3.3 Feature importance graph based on SHAP values:

This figure represents the SHAP workflow, highlighting words that contribute positively (in green/red) toward a specific outcome. It validates that the system highlights essential legal entities while assigning zero-weight to grammatical noise.




Fig 6: Feature importance graph based on SHAP values

3.4 ROC curve of the InLegalBERT classifier:

The ROC curve, achieving an **AUC score of 0.91**, validates that our fine-tuned model holds superior threshold discrimination compared to a random classifier, ensuring high-confidence predictions for practitioners.




Fig 7: ROC curve of the InLegalBERT classifier

3.5 Feature correlation heatmap:

This heatmap evaluates the interactions between different legal charges and verdicts. It reveals significant positive and negative relationships (e.g., between murder charges and appellate verdicts), proving that our model captures the joint occurrence trends within Indian legal texts.



Fig 8: Feature correlation heatmap

3.6 Algorithm runtime comparison:

This chart compares the mean inference time per document for InLegalBERT versus baseline models. It highlights the trade-off between predictive accuracy and computational speed, confirming the system is commercially ready for high-throughput analysis.




Fig 9: Algorithm runtime comparison

3.7 Prediction Confidence Distribution (PCD):

The PCD graph visualizes the model's discriminative confidence across the test set. By plotting the density of favorable versus unfavorable classifications, it shows that the model is highly certain, reducing the risk of neutral or uncertain decision-making.




Fig 10: Prediction Confidence Distribution (PCD)

3.8 The platform login interface for secure access:

This figure displays the encrypted authentication system, which serves as the principal entrance to the predictive context. It ensures that sensitive legal analysis and AI-generated insights are restricted to authorized practitioners.





Fig 11: The platform login interface for secure access

3.9 The Secure Case History dashboard:

The case history interface functions as a professional audit log. It integrates with the cloud backend to provide practitioners with a permanent, encrypted record of AI-driven insights for comparative analysis and case strategy.




created_at	filename	predicted_jurisdiction	predicted_outcome	outcome_confidence
2026-03-25T15:10:53.073676+00:00	10718_2011_7_1502_53219_Judgement_15-May-2024.pdf	Criminal Law	Allowed / Accepted	96.9868
2026-03-25T10:10:17.239154+00:00	perfect_test_case.txt	Criminal Law	Dismissed / Rejected	53.8859
2026-03-25T10:04:57.902965+00:00	ordjud.pdf	Criminal Law	Allowed / Accepted	91.8982
2026-03-20T20:06:01.95502+00:00	perfect_test_case.txt	Criminal Law	Dismissed / Rejected	53.8859
2026-03-20T19:29:33.019161+00:00	ordjud.pdf	Criminal Law	Allowed / Accepted	91.8982
2026-03-16T19:16:31.4429+00:00	perfect_test_case.txt	Criminal Law	Dismissed / Rejected	53.9647

Fig 12: The Secure Case History dashboard

3.10 The detailed case analysis dashboard:

This dashboard acts as the primary workspace, integrating prediction overviews with in-depth structural analysis. It provides multi-level perspectives on document legal factors, empowering practitioners to perform quick audits of the model's reasoning.



created_at	filename	predicted_jurisdiction	predicted_outcome	outcome_confidence
2026-03-25T15:10:53.073676+00:00	10718_2011_7_1502_53219_Judgement_15-May-2024.pdf	Criminal Law	Allowed / Accepted	96.9868
2026-03-25T10:10:17.239154+00:00	perfect_test_case.txt	Criminal Law	Dismissed / Rejected	53.8859
2026-03-25T10:04:57.902965+00:00	ordjud.pdf	Criminal Law	Allowed / Accepted	91.8982
2026-03-20T20:06:01.95502+00:00	perfect_test_case.txt	Criminal Law	Dismissed / Rejected	53.8859
2026-03-20T19:29:33.019161+00:00	ordjud.pdf	Criminal Law	Allowed / Accepted	91.8982
2026-03-16T19:16:31.4429+00:00	perfect_test_case.txt	Criminal Law	Dismissed / Rejected	53.9647

Predicted Jurisdiction
Criminal Law
Confidence: 96.9%

Predicted Appellate Outcome
Allowed / Accepted
Confidence: 97.2%

Document Structure Analysis (AI Extracted)

Extracted Facts: Selected Arguments: Full Mapped Prelude

Notice how IPC references have been automatically converted to BNS equivalents in brackets.
2024 INSC 422 REPORTABLE IN THE SUPREME COURT OF INDIA CRIMINAL APPELLATE JURISDICTION
CRIMINAL APPEAL NO. 2281 OF 2011
BALASOBA S/O RAMDAS HOLNE
APPELLANT(S)
VERSUS
STATE OF MADHARASHASTRA
2.
2.1.
It may be mentioned that by the judgment and order dated 22.07.2008, the trial court had convicted the appellant for committing an offence punishable under Section 303 of BNS) read with Section 34 of the [Bharatiya Nyaya Sanhita (BNS)], 1961) (BNS) and sentenced to suffer the imprisonment and to pay a fine of Rs. 20,000/- with a default stipulation. The appeal filed by the appellant against the aforesaid conviction and sentence was dismissed by the High Court. 3.
The prosecution case is brief in that wife of the appellant Rekha was a police constable and lived in the police colony at Ambikagali. Her husband i.e. the appellant was serving in the army. He had come home on leave. 3.1.
On 22.07.2002, at about 08.30 PM, Rekha had sustained burn injuries in the quarter where she was residing. According to the prosecution, she was subjected to cruelty by her husband Rajendra and brother-in-law Suresh. She was also subjected to sustained cruelty at the hands of her other in-laws including father-in-law, mother-in-law and sister-in-law. On the fateful day, Rekha was beaten by her husband Rajendra and brother-in-law Suresh. They tied her hands with a garbha and her feet by a towel. Then the husband gagged her back. 3.1.1.

Fig 13: The detailed case analysis dashboard

4. Conclusions

In this project, we successfully engineered a production-grade, full-stack machine learning system to predict the outcomes of Indian legal cases. By addressing complex challenges such as the legal evolution from the IPC to the BNS via dynamic middleware, and overcoming standard transformer token limits through structure-aware chunking, the architecture has proven to be highly robust for modern jurisprudence.

Furthermore, the integration of a secure Supabase backend with an XAI-qualified Streamlit frontend effectively transforms a theoretical AI model into a practical, transparent decision-support tool. Our system establishes that Artificial Intelligence holds significant potential for transforming legal research and case strategy, reducing the time required for manual analysis while enhancing consistency in judicial forecasting.

Future Work

There are several exciting opportunities to enhance the **Legal AI Hub** in future iterations:

- **Generative Summarization:** Integrating larger LLMs to provide automated, generative case summarization alongside predictive outcomes.
- **Architecture Scaling:** Transitioning the frontend from Streamlit to a fully modified React.js or Next.js architecture to allow for greater UI/UX flexibility and continuous mobile support for legal professionals on the go.
- **Regional Legal Expansion:** Increasing the dynamic JSON vocabularies to include regional state laws, which would considerably extend the system's general efficacy across the diverse Indian legal ecosystem.

In summary, the successful deployment of this platform validates that domain-specific transformer models, when combined with structured data engineering, provide a reliable pathway for automated judicial analysis. This project serves as a foundational step toward more transparent and efficient legal technology in India.

Research Papers Published/Certificates

- **Raghuwanshi, N., et al.** "AI for Legal Case Outcome Prediction." [Manuscript submitted for publication/presented at Nutan College of Engineering and Research].
- **Certificate of Completion:** Project Development – AI for Legal Case Outcome Prediction, [2025-26].
- **Participation Certificate:** [Insert name of Hackathon/Conference], [Year].

References

1. Ben-Ari, R. "Computational law: Towards a new paradigm." *Proceedings of the 2012 IEEE International Conference on Systems, Man, and Cybernetics (SMC)*, 2012, pp. 2480-2485.
2. Ashley, K. D. *Artificial Intelligence and Legal Analytics: New Tools for Law Practice in the Digital Age*. Cambridge University Press, 2017.
3. Vaswani, A., et al. "Attention Is All You Need." *Advances in Neural Information Processing Systems*, 2017, pp. 5998-6008.
4. Lundberg, S. M., and Lee, S. I. "A Unified Approach to Interpreting Model Predictions." *Advances in Neural Information Processing Systems*, 2017, pp. 4765-4774.
5. Al-Kassab, K., et al. "Legal Judgment Prediction: A Comparative Study of Machine Learning Approaches." *Proceedings of the International Conference on Artificial Intelligence and Law (ICAIL)*, 2021.
6. Devlin, J., et al. "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." *Proceedings of NAACL-HLT*, 2019, pp. 4171-4186.
7. Chalkidis, P., et al. "Legal-BERT: The Technical Challenges of Adapting BERT for Legal NLP." *Proceedings of the 28th International Conference on Legal Knowledge and Information Systems (JURIX)*, 2019.
8. Malik, V., et al. "ILDC for CJPE: Indian Legal Documents Corpus for Court Judgment Prediction and Explanation." *Proceedings of ACL-IJCNLP*, 2021, pp. 4046-4062.
9. Nigam, S. K., et al. "NyayaAnumana & INLegalLlama: The Largest Indian Legal Judgment Prediction Dataset." *arXiv preprint arXiv:2412.08385*, 2024.
10. Nigam, S. K., et al. "TathyaNyaya and FactLegalLlama: Advancing Factual Judgment Prediction." *arXiv preprint arXiv:2504.04737*, 2025.
11. Aletras, N., et al. "Predicting judicial decisions of the European Court of Human Rights." *PeerJ Computer Science*, 2016.
12. Zhong, Z., et al. "Legal Judgment Prediction via Topological Learning." *Proceedings of EMNLP*, 2018.
13. Feng, Y., et al. "Legal Judgment Prediction via Event Extraction." *Proceedings of COLING*, 2022.
14. Šavelka, J., and Ashley, K. "Leveraging Large Language Models for Legal Text Prediction." *Artificial Intelligence and Law*, 2023.
15. Katsikouli, P., et al. "Machine Learning and Asylum Adjudications: From Analysis of Variations to Outcome Predictions." *2022 IEEE International Conference on Big Data*, 2022.

16. Kmainasi, M. B., et al. "Can Large Language Models Predict the Outcome of Judicial Decisions?" *arXiv preprint arXiv:2501.09768*, 2025.
17. Ariai, F., and Demartini, G. "Natural Language Processing for the Legal Domain: A Survey." *arXiv preprint arXiv:2410.21306*, 2024.
18. Sahoo, S., et al. "A Comprehensive Analysis of Indian Legal Documents Summarization Techniques." *Neural Computing and Applications*, 2023.
19. Jaegers, R. "A Novel Approach to Legal Text Summarization with Deep Reinforcement Learning." *Journal of Political Science and International Relations Studies*, 2024.
20. de Souza, P., et al. "Topic Modelling of Legal Documents via LEGAL-BERT." *2022 IEEE International Conference on Big Data*, 2022.